Data Engineering Bootcamp

First Delivery Report

Autor: Alejandra Elizabeth Moreno Morales

Task description

Based on the self-study material, recorded and live session, and mentorship covered until this deliverable, we suggest you perform the following:
Take as reference Terraform reference, identify and select the corresponding terraform blocks to build your own Airflow Cluster.
Airflow Cluster must be built with GKS in Google or EKS in AWS.
In case of some difficulties, take advantage of templates provided by Wizeline to build and start your Airflow Cluster.
Take your notes about any blocker and your lessons learned to be discussed during Q&A and Mentoring sessions.

Target

Implement configuration for run an Airflow cluster in a Google Kubernetes Services (GKS) of Google Cloud Platform (GCP) using Terraform as tool to declare the configuration files.

Prerequisites

- Terraform configuration

To run configuration files with the declaration of the infrastructure as code.

- GCP account

To deploy resources requested
Be aware of:
-Configure a Service account to be used on behalf of terraform

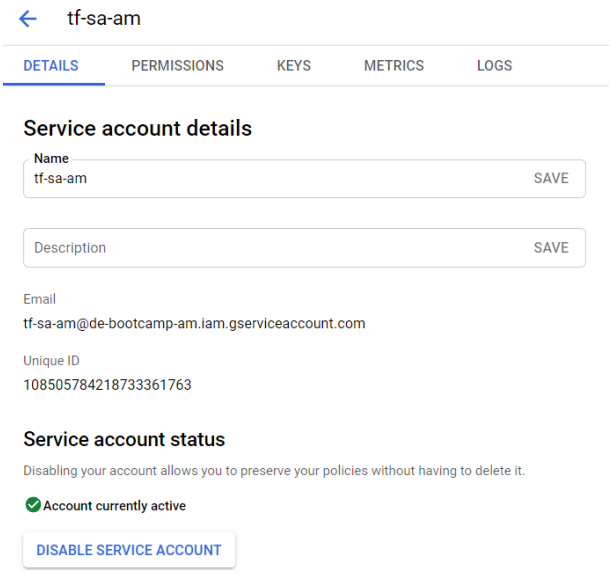

-Enable Compute Engine API and Kubernetes Engine API in order to Terraform could work on this configuration.

- GCloud SDK tool configuration

To be used as an access for GCP using your user account credentials and therefore Terraform be allowed to provision resources on GCloud.

- Kubectl tool configuration

To be able to control Kubernetes clusters

- Helm3

To manage the Kubernetes application.

Terraform configuration files

- main.tf. Declare provisions needed. A VPC and subnet will be created. Also deploy a 2-node separately managed node pool GKE cluster. And SQL resource running PostgreSQL 12.
- output.tf. Define outputs values after creation.
- provider.tf. Configures the specified provider to Terraform uses to create and manage your resources.
- terraform.tfvars. Template used to set the values for variables
- variables.tf. Declare the variables name, description and default value to be used in the project.

Implementation

1. Create and customize the Terraform files to deploy the infrastructure needed. See Terraform configuration files
2. Initialize your Terraform workspace to install the plugins needed to manage the infrastructure

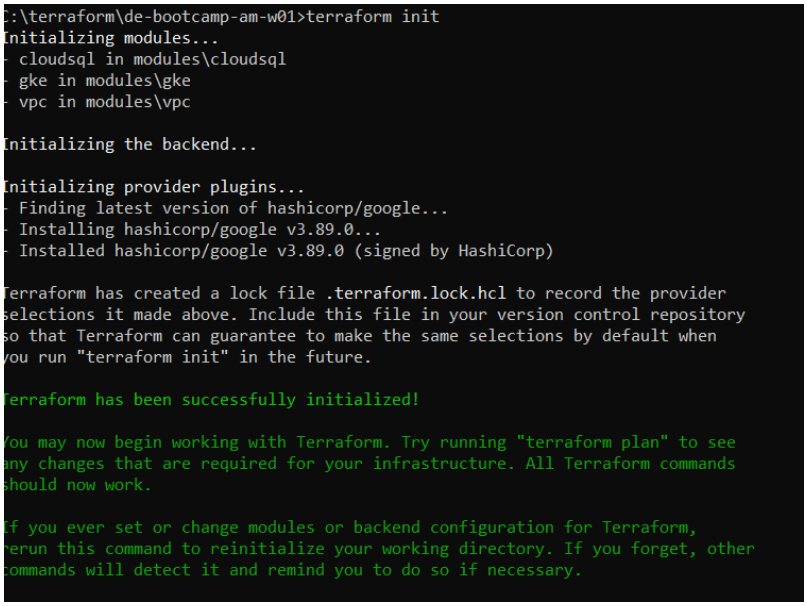

Create EKS cluster

3. Run Terraform apply and review the planned actions. Your terminal output should indicate the plan is running and what resources will be created. Confirm the apply.

```
C:\terraform\de-bootcamp-am-w01>terraform apply --var-file=terraform.tfvars
module.vpc.google_compute_network.main-vpc: Refreshing state... [id=projects/de-bootcamp-am/global/networks/de-bootcamp-am-vpc]
module.cloudsql.google_sql_database_instance.sql_instance: Refreshing state... [id=data-bootcamp-am-1]
module.vpc.google_compute_subnetwork.private_subnets[0]: Refreshing state... [id=projects/de-bootcamp-am/regions/us-central1/subnetworks/private-0-private-subnet]
module.vpc.google_compute_subnetwork.private_subnets[1]: Refreshing state... [id=projects/de-bootcamp-am/regions/us-central1/subnetworks/private-1-private-subnet]
module.vpc.google_compute_subnetwork.public_subnets[0]: Refreshing state... [id=projects/de-bootcamp-am/regions/us-central1/subnetworks/public-0-public-subnet]
module.vpc.google_compute_subnetwork.private_subnets[2]: Refreshing state... [id=projects/de-bootcamp-am/regions/us-central1/subnetworks/private-2-private-subnet]
module.vpc.google_compute_subnetwork.public_subnets[2]: Refreshing state... [id=projects/de-bootcamp-am/regions/us-central1/subnetworks/public-2-public-subnet]
module.vpc.google_compute_subnetwork.public_subnets[1]: Refreshing state... [id=projects/de-bootcamp-am/regions/us-central1/subnetworks/public-1-public-subnet]
module.cloudsql.google_sql_database.database: Refreshing state... [id=projects/de-bootcamp-am/instances/data-bootcamp-am-1/databases/dbname-am]
module.gke.google_container_cluster.primary: Refreshing state... [id=projects/de-bootcamp-am/locations/us-central1-a/clusters/airflow-gke-data-bootcamp]
module.gke.google_container_node_pool.primary_nodes: Refreshing state... [id=projects/de-bootcamp-am/locations/us-central1-a/clusters/airflow-gke-data-bootcamp/nod
ePools/airflow-gke-data-bootcamp-node-pool]
```

4. Upon successful application, your terminal prints the outputs and save in terraform.tfstate

terraform.tfstate

C: > terraform > de-bootcamp-am-w01 > terraform.tfstate


```
{
  "version": 4,
  "terraform_version": "1.0.9",
  "serial": 15,
  "lineage": "dfa33958-400e-ce7b-0c2b-44494d64f41b",
  "outputs": {
    "kubernetes_cluster_host": {
      "value": "104.154.26.35",
      "type": "string"
    },
    "kubernetes_cluster_name": {
      "value": "airflow-gke-data-bootcamp",
      "type": "string"
    },
    "location": {
      "value": "us-central1-a",
      "type": "string"
    },
    "project_id": {
      "value": "de-bootcamp-am",
      "type": "string"
    },
    "region": {
      "value": "us-central1",
      "type": "string"
    }
  },
  "resources": [
    {
      "module": "module.cloudsql",
      "mode": "managed",
      "type": "google_sql_database",
      "name": "database",
      "provider": "provider[\"registry.terraform.io/hashicorp/google\"]",
      "instances": [
        {
```


5. Review resources created in GCP console

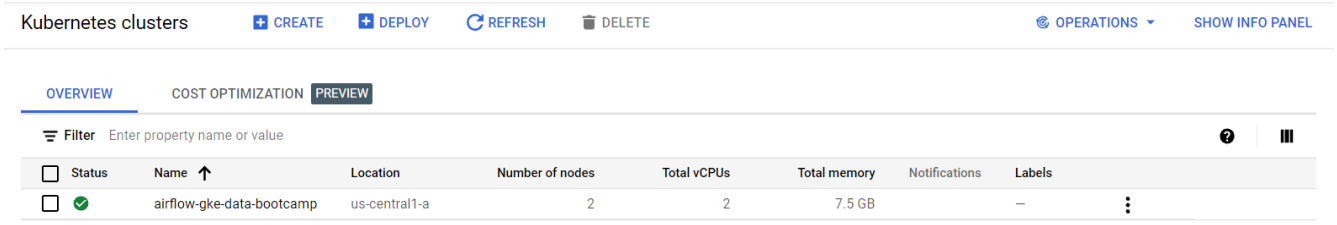

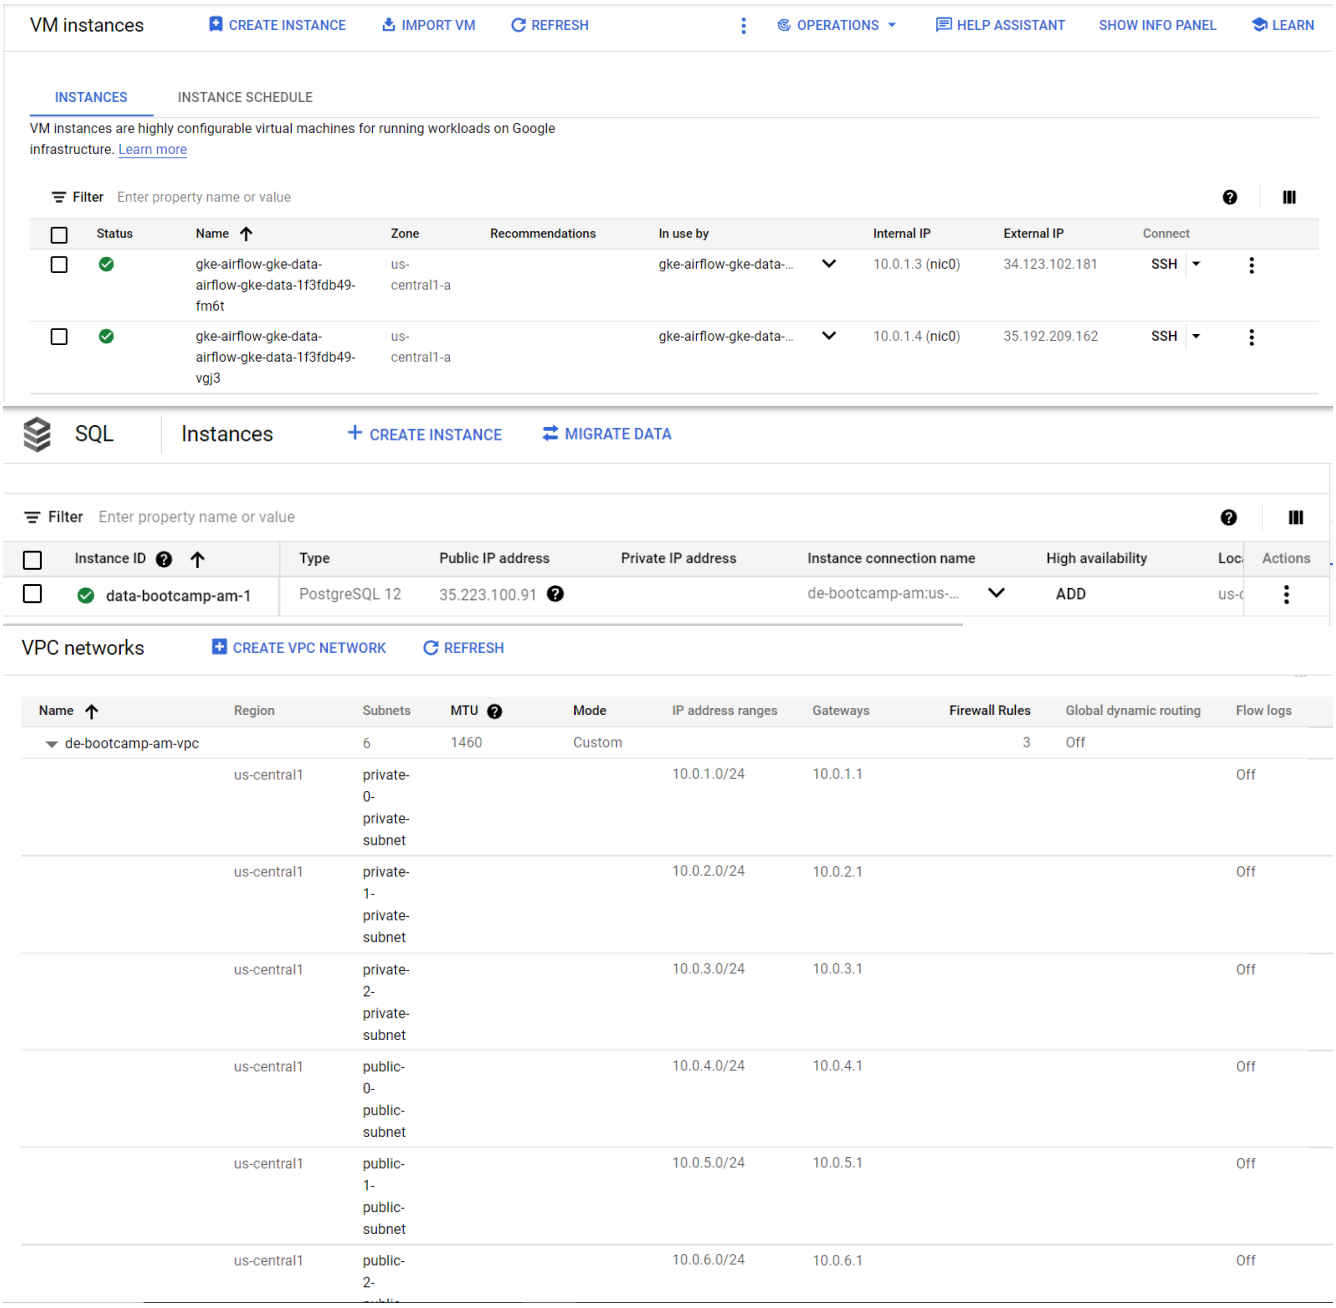

6. Once that the cluster is created, set the kubectl context

```
C:\terraform\de-bootcamp-am-w01>gcloud container clusters get-credentials airflow-gke-data-bootcamp --zone=us-central1-a
Fetching cluster endpoint and auth data.
kubeconfig entry generated for airflow-gke-data-bootcamp.
```

Create NFS Service

7. Create a namespace for the nsf service

```
C:\terraform\de-bootcamp-am-w01>kubectl create namespace nfs
namespace/nfs created
```

8. Create the nfs server

```
C:\terraform\de-bootcamp-am-w01>kubectl -n nfs apply -f nfs/nfs-server.yaml
persistentvolumeclaim/nfs-pvc created
deployment.apps/nfs-server created
service/nfs-server created
```

Create Storage

9. Create a namespace for storage deployment

```
C:\terraform\de-bootcamp-am-w01>kubectl create namespace storage
namespace/storage created
```

10. Add the chart for the nfs-provisioner

```
C:\terraform\de-bootcamp-am-w01>helm repo add nfs-subdir-external-provisioner https://kubernetes-sigs.github.io/nfs-subdir-external-provisioner/
"nfs-subdir-external-provisioner" has been added to your repositories
```

11. Install nfs-external-provisioner

```
C:\terraform\de-bootcamp-am-w01>helm install nfs-subdir-external-provisioner nfs-subdir-external-provisioner/nfs-subdir-external-provisioner --namespace storage --set nfs.server=10.7.251.128 --set nfs.path=/
NAME: nfs-subdir-external-provisioner
LAST DEPLOYED: Sun Oct 24 21:40:00 2021
NAMESPACE: storage
STATUS: deployed
REVISION: 1
TEST SUITE: None
```

Create Airflow

12. Create a namespace for airflow deployment

```
C:\terraform\de-bootcamp-am-w01>kubectl create namespace airflow
namespace/airflow created
```

13. Add the chart repository

```
C:\terraform\de-bootcamp-am-w01>helm repo add apache-airflow https://airflow.apache.org
"apache-airflow" has been added to your repositories
```

14. Update airflow-values.yaml file with the project values

airflow-values.yaml

C: > terraform > de-bootcamp-am-w01 > kubernetes > ! airflow-values.yaml

```yaml
# Git sync
dags:
  persistence:
    # Enable persistent volume for storing dags
    enabled: true
    # Volume size for dags
    size: 1Gi
    # If using a custom storageClass, pass name here
    storageClassName: nfs-client
    # access mode of the persistent volume
    accessMode: ReadWriteMany

  gitSync:
    enabled: true

    # git repo clone url
    # ssh examples ssh://git@github.com/apache/airflow.git
    # git@github.com:apache/airflow.git
    # https example: https://github.com/apache/airflow.git
    repo: https://github.com/aleksmoreno2/Airflow-Templates.git
    branch: main
    rev: HEAD
    depth: 1
    # the number of consecutive failures allowed before aborting
    maxFailures: 0
    # subpath within the repo where dags are located
    # should be "" if dags are at repo root
    subPath: ""
```

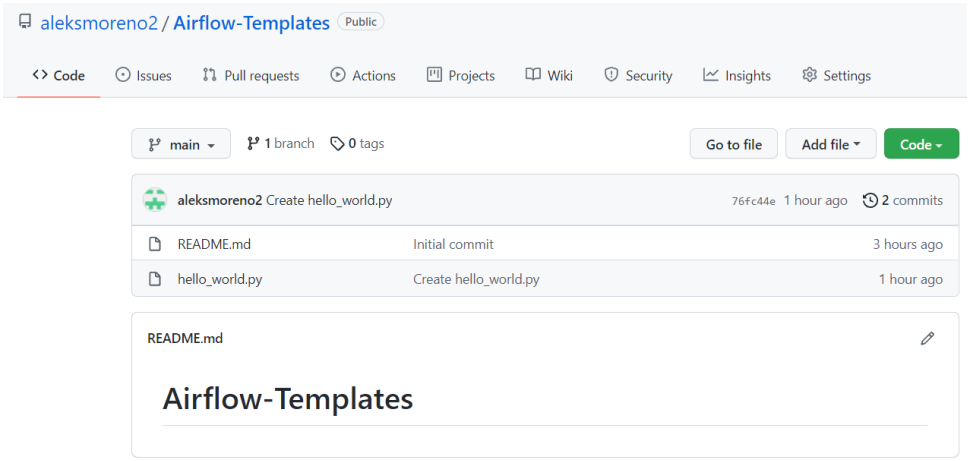

15. Install the airflow chart from the repository

```
C:\terraform\de-bootcamp-am-w01>cd kubernetes

C:\terraform\de-bootcamp-am-w01\kubernetes>helm install airflow -f airflow-values.yaml apache-airflow/airflow --namespace airflow
NAME: airflow
LAST DEPLOYED: Sun Oct 24 19:24:57 2021
NAMESPACE: airflow
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
Thank you for installing Apache Airflow 2.1.4!

Your release is named airflow.
You can now access your dashboard(s) by executing the following command(s) and visiting the corresponding port at localhost in your browser:

Airflow Webserver:     kubectl port-forward svc/airflow-webserver 8080:8080 --namespace airflow
Flower dashboard:      kubectl port-forward svc/airflow-flower 5555:5555 --namespace airflow
Default Webserver (Airflow UI) Login credentials:
    username: admin
    password: admin
Default Postgres connection credentials:
    username: postgres
    password: postgres
    port: 5432

You can get Fernet Key value by running the following:

    echo Fernet Key: $(kubectl get secret --namespace airflow airflow-fernet-key -o jsonpath="{.data.fernet-key}" | base64 --decode)

###########################################################
#  WARNING: You should set a static webserver secret key  #
###########################################################

You are using a dynamically generated webserver secret key, which can lead to
unnecessary restarts of your Airflow components.

Information on how to set a static webserver secret key can be found here:
https://airflow.apache.org/docs/helm-chart/stable/production-guide.html#webserver-secret-key
```

16. Verify that our pods are up and running

```
C:\terraform\de-bootcamp-am-w01>kubectl get pods -n airflow
NAME                                 READY   STATUS    RESTARTS   AGE
airflow-flower-6c6b7f5d68-z275g      1/1     Running   1          11m
airflow-postgresql-0                 1/1     Running   0          11m
airflow-redis-0                      1/1     Running   0          11m
airflow-scheduler-56fbb444-gc6nm     3/3     Running   0          11m
airflow-statsd-84f4f9898-hxg46       1/1     Running   0          11m
airflow-webserver-66f7788c78-f4jjx   1/1     Running   0          11m
airflow-worker-0                     2/2     Running   0          7m9s
```

17. Accessing to Airflow dashboard

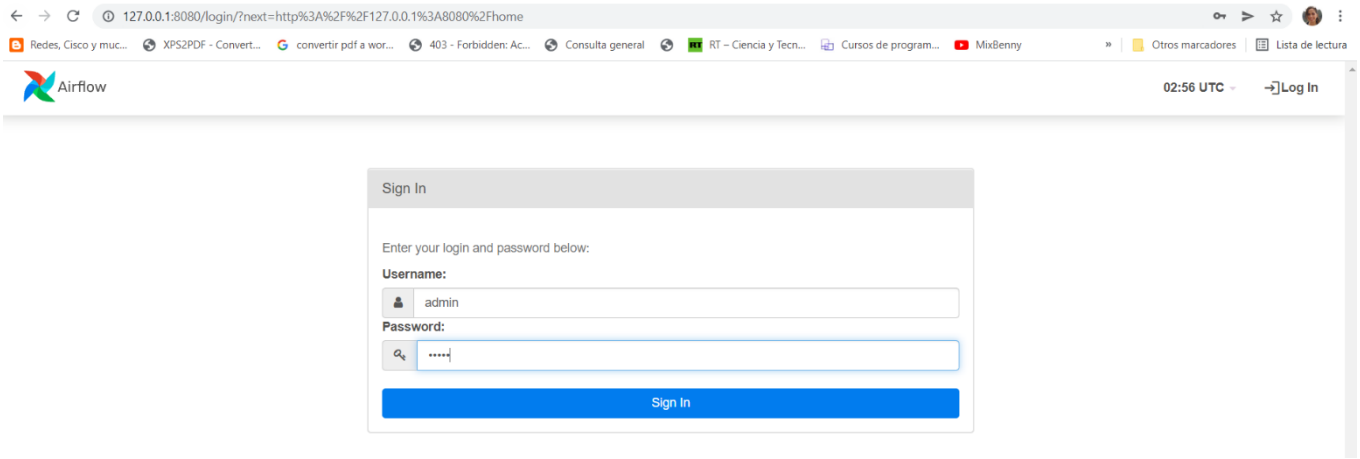

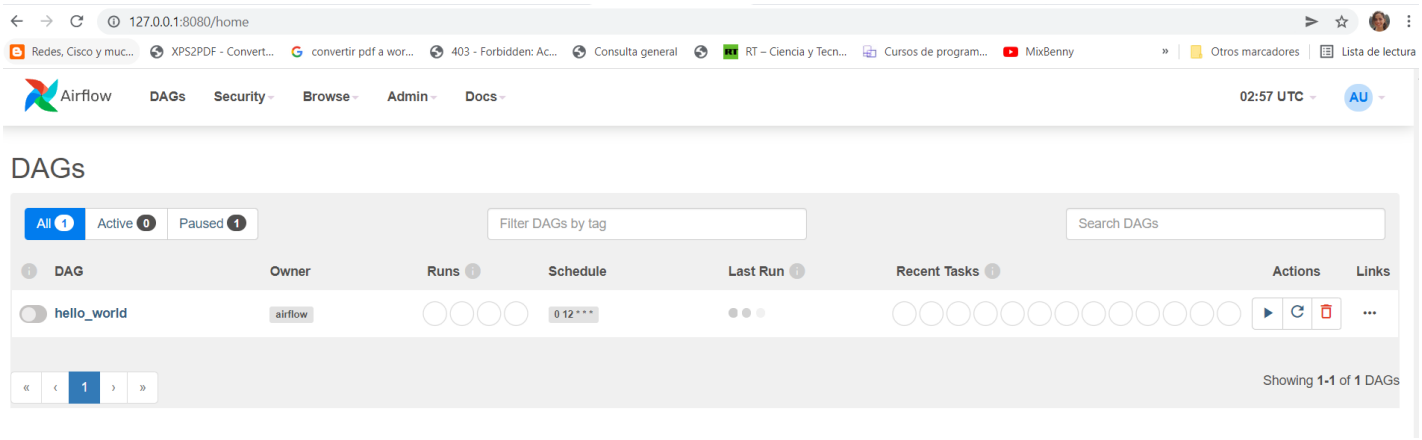
127.0.0.1:8080/home
Redes, Cisco y muc...
XPS2PDF - Convert...
convertir pdf a wor...
403 - Forbidden: Ac...
Consulta general
RT – Ciencia y Tecn...
Cursos de program...
MixBenny
Otros marcadores
Lista de lectura
Airflow
DAGs
Security
Browse
Admin
Docs
02:57 UTC
AU
DAGs
All 1
Active 0
Paused 1
Filter DAGs by tag
Search DAGs
DAG
Owner
Runs
Schedule
Last Run
Recent Tasks
Actions
Links
hello_world
airflow
0 12 * * *
Showing 1-1 of 1 DAGs